



Business Informatics Group

**Vienna University of Technology**

# Web Engineering

Introduction to JPA and Hibernate



**Philipp Liegl**

**Business Informatics Group**

Institute of Software Technology and Interactive Systems  
Vienna University of Technology

Favoritenstraße 9-11/188-3, 1040 Vienna, Austria  
phone: +43 (1) 58801-18804 (secretary), fax: +43 (1) 58801-18896  
office@big.tuwien.ac.at, www.big.tuwien.ac.at

# Outline of today's talk

---

- JDBC
- JPA/Hibernate
  - Relationships
  - Persistence Context/Persistence Unit
  - Entity Manager
  - JPQL
  - Hibernate Criteria API

Accompanying examples

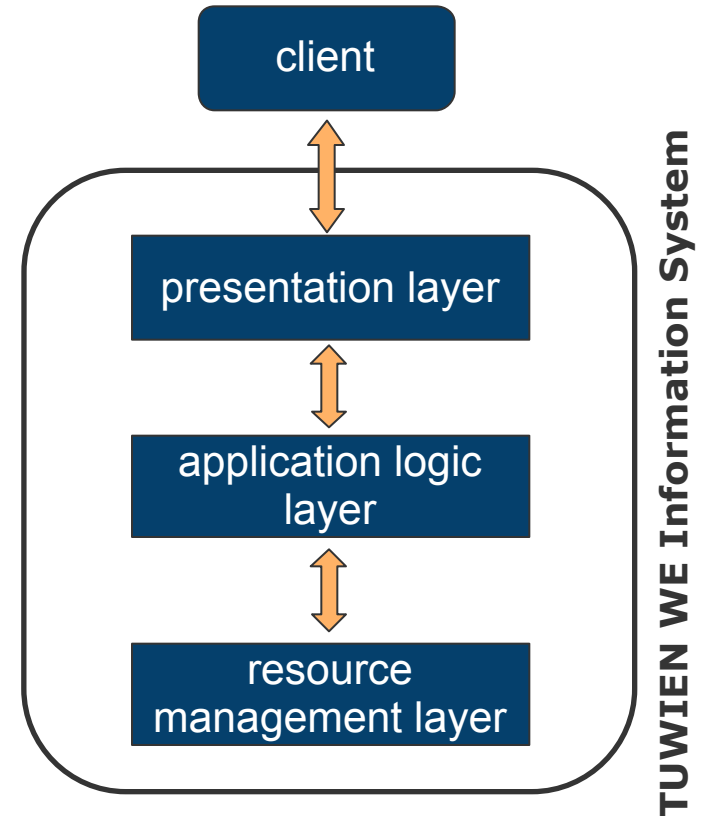
<https://github.com/pliegl/we2014/tree/master/jpa-sample>

# Motivation

## N-Tier Architectures

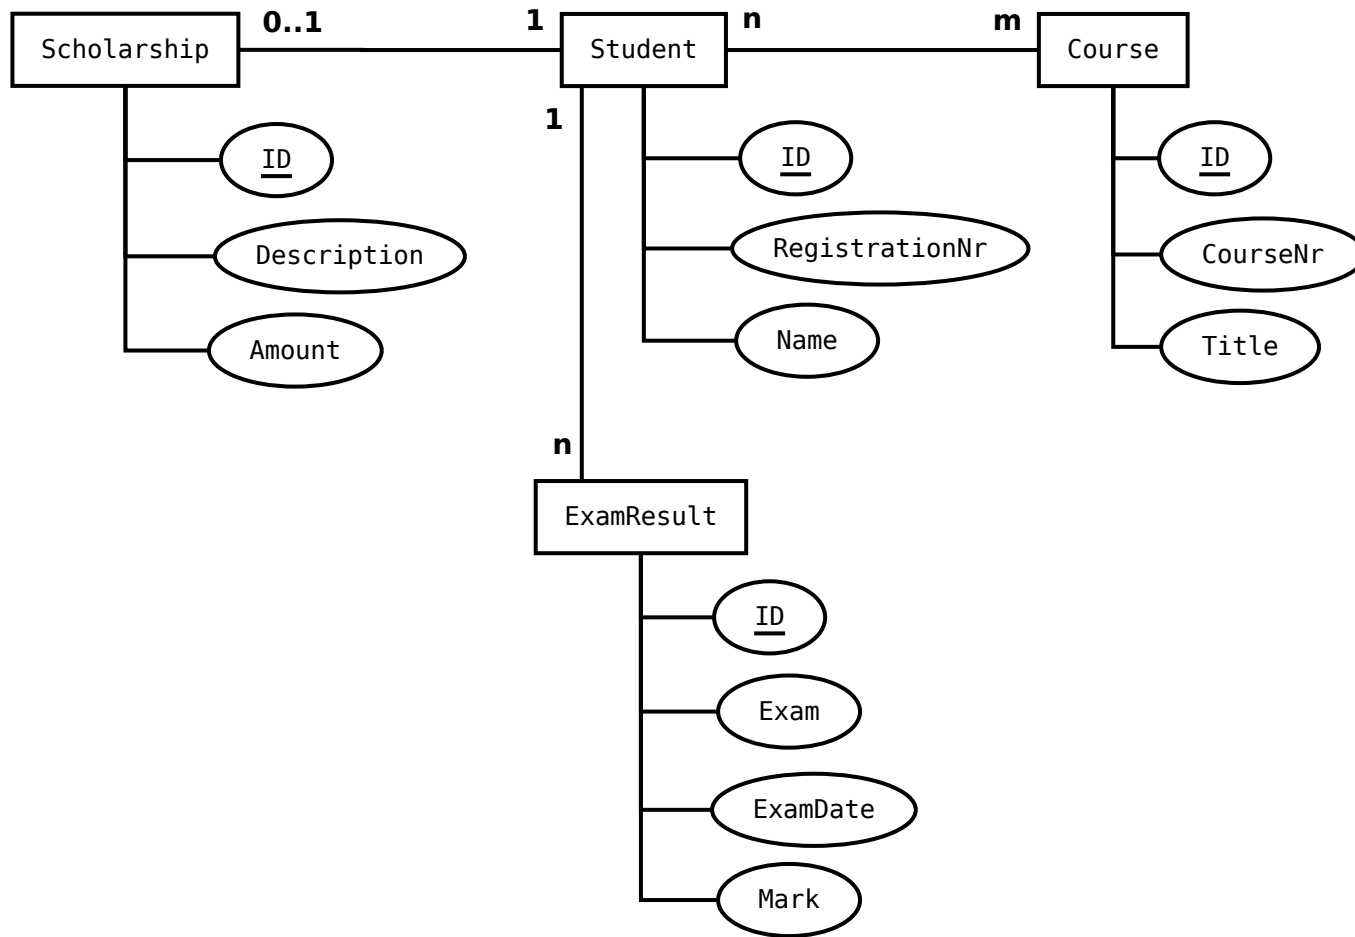
### ■ Layers of an information system

- Presentation layer
  - Communication interface to external entities
  - “View” in the model-view-controller
- Application logic layer (service layer)
  - Implements operations requested by clients through the presentation layer
  - Represents the “business logic”
- Resource management layer (persistence layer)
  - Deals with different data sources of an information system
  - Responsible for storing and retrieving data




# Motivation

Accompanying model



# Motivation

## Traditional persistence with JDBC



# Motivation

## JDBC - Java Database Connectivity

---

- Used to access relational databases from Java programs
- First version released 1996
- Ability to
  - Establish a connection to a database
  - Execute an SQL statement and return results
  - Create parameterized queries
  - Manage database transactions
- Basic Steps
  - Load driver or obtain an already defined data source
  - Establish connection using a JDBC URL
  - Create an **SQL statement** and execute SQL statement
  - If present, process results present in **result sets**
  - Close database resources
  - Commit or rollback transaction, if necessary



# JDBC

## Insert an entry

```
Connection conn = null;
PreparedStatement stmt = null;

try {
    conn = connection();
    stmt = conn.prepareStatement( "INSERT INTO student VALUES(?, ?, ?)" );
    stmt.setInt( 1, student.getId() );
    stmt.setString( 2, student.getMatrNr() );
    stmt.setString( 3, student.getName() );
    stmt.executeUpdate();
    stmt.close();
} catch (Exception e) {
    e.printStackTrace();
} finally {
    if (stmt != null) {
        stmt.close();
    }
    if (conn != null) {
        conn.close();
    }
}
```

# JDBC

## Retrieve an entry

```
Connection conn = null;
PreparedStatement stmt = null;
ResultSet rs = null;
try {
    conn = connection();
    stmt = conn.prepareStatement( "SELECT id, matnr, name FROM student
                                  WHERE id=?" );

    stmt.setInt( 1, id );
    rs = stmt.executeQuery();
    rs.next();

    Student student = new Student();
    student.setId( rs.getInt( 1 ) );
    student.setMatrNr( rs.getString( 2 ) );
    student.setName( rs.getString( 3 ) );

    rs.close();
    stmt.close();
    return student;
} catch (Exception e) {
    e.printStackTrace();
} finally {
    ...
}
```



# Java Persistence API (JPA)

## Introduction

---

- Specification for the management of persistence and object/relational mapping with Java
  - Persistence: Data objects shall outlive the JVM app
- Objective: provide an object/relational mapping facility for Java developers using a Java domain model and a relational database
  - Map Java POJOs to relational databases (which are one type of persistence)
- Standardized under the Java Community Process Program with contributions from Hibernate, TopLink, JDO, and the EJB community
- Hibernate: JPA implementation with additional “native” features, e.g.,
  - HQL (Hibernate Query Language) - similar to JPQL, but with some extensions
  - Criteria API



# Object Relational Mapping

## Reasons for using ORM

---

- In an application we want to focus on **business concepts**, not on the relational database structure
- Abstract from the “by-hand” communication with the DB (e.g., via JDBC)
- Allow for an automatic synchronization between Java Objects and the underlying database
- Portability
  - Mostly DB independent (with the exception of some types of features, such as identifier generation)
  - Query abstractions using e.g. JPQL - the vendor specific SQL is auto-generated
- Performance
  - Object and query caching is automatically done by the ORM



# Persistent Entities


## Basics

---

- Are POJOs (Plain Old Java Objects)
- Lightweight persistent domain object
- Typically represent a table in a relational database
- Each entity instance corresponds to one row in that table
- Have a persistent identity
- May have both, persistent and transient (non-persistent) state
  - Simple types (primitive data types, wrappers, enums)
  - Composite types (e.g., Address)
  - Non-persistent state (using identifier `transient` or `@Transient` annotation)



# Persistence with Hibernate



# Simple Mapping

Enhance Java domain classes with JPA annotations

```
@Entity
public class ExamResult {

    @Id
    private Long id;

    @Column(name = "pruefungsDatum")
    @Temporal(TemporalType.DATE)
    private Date examDate;

    private int mark;

    @Transient
    private String examLocation;

    //Getter and setters omitted
}
```

## ExamResult

id	pruefungsDatum	mark
----	----------------	------

### Important annotations

@Entity	Specifies that the class is an entity
@Id	Specifies the primary key of an entity
@Temporal	Must be specified for fields of type java.util.Date and java.util.Calendar
@TemporalType	Type used to indicate a specific mapping of java.util.Date or java.util.Calendar. Allowed values: <ul style="list-style-type: none"><li>- DATE</li><li>- TIME</li><li>- TIMESTAMP</li></ul>
@Transient	Specifies that the field is not persistent



# Simple Mapping

## Inheritance

```
@MappedSuperclass
public class BaseEntity {

    @Id
    @GeneratedValue
    protected Long id;

    public Long getId() {
        return id;
    }
}
```

### Important annotations

<b>@MappedSuperclass</b>	Designates a class whose mapping information is applied to the entities that inherit from it. A mapped superclass has no separate table defined for it.
<b>@GeneratedValue</b>	The <code>GeneratedValue</code> annotation may be applied to a primary key property or field of an entity or mapped superclass in conjunction with the <code>Id</code> annotation.

```
@Entity
public class ExamResult extends BaseEntity {

    @Column(name = "pruefungsDatum")
    @Temporal(TemporalType.DATE)
    private Date examDate;

    private int mark;

    @Transient
    private String examLocation;

    //Getter and setters omitted

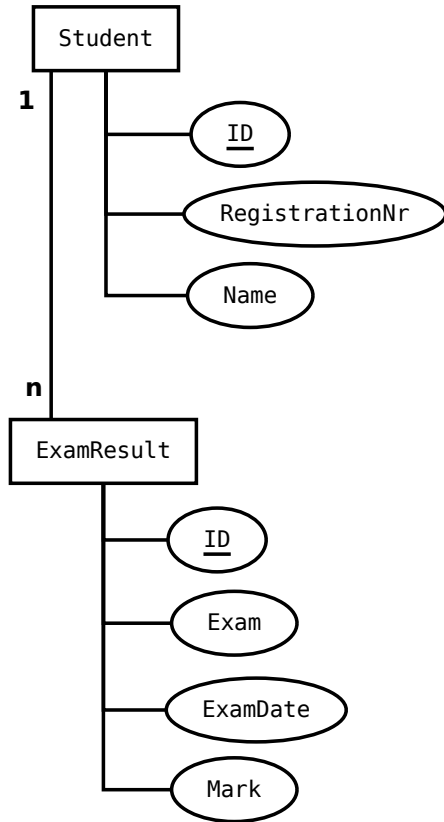
}
```



ExamResult		
id	pruefungsDatum	mark



# Object-oriented vs. SQL



OO: **Student owns the ExamResults**  
**Usually:** no ExamResult without a student

```
public class Student extends BaseEntity {  
  
    private String registrationNumber;  
    private String name;  
    private List<ExamResult> examResults;  
    ...  
}
```

Does not exist in the DB, but is simulated using an SQL query

```
public class ExamResult extends BaseEntity {  
  
    private Date examDate;  
    private String exam;  
    private int mark;  
    private Student student;  
    ...  
}
```

SQL:

- ExamResult contains a foreign key to the Student it belongs to
- The ExamResult owns (contains) the connection
- This is opposite to the OO perspective

# Entity relationships

---

**One-to-one, one-to-many, many-to-many, many-to-one** relationships among entities

- bi-directional or uni-directional
- Support for different Collection types, e.g., List, Set, Map, etc.

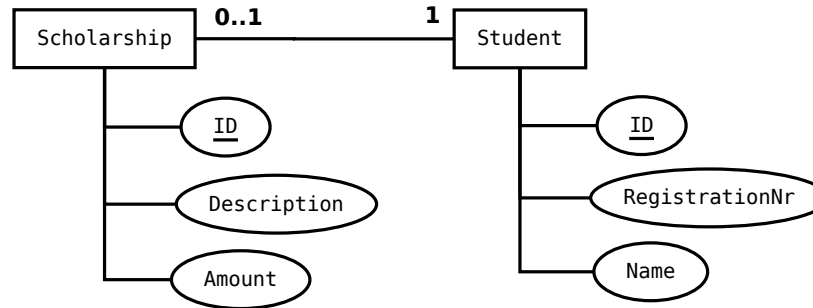
Need to specify the **owning** side in relationships

- Owning side table has the foreign key
- OneToOne relationship - the side where the foreign key is specified
- OneToMany, ManyToOne - the “many” side



# Relationship mapping

Example using a unidirectional mapping

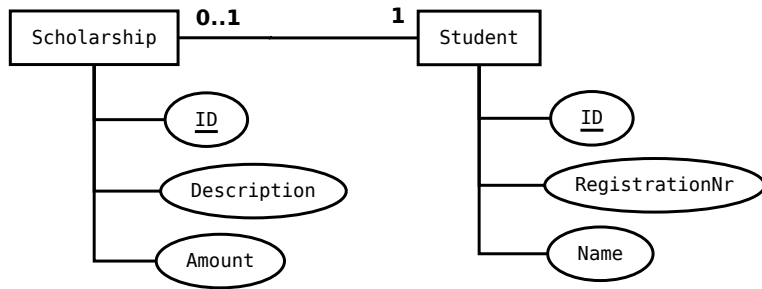


## Four different options:

1. Using an embedded table, where Scholarship is the embedded table. (see example)
2. Scholarship and Student are separate tables. The primary key of Scholarship has a foreign key constraint on the primary key of the “owning” Student (using `@PrimaryKeyJoinColumn` annotation)
3. Scholarship and Student are separate tables. Student holds a foreign key which references the primary key of Scholarship. The foreign key has a unique constraint.
4. Scholarship and Student are separate tables. Scholarship holds a foreign key which references the primary key of Student. The foreign key has a unique constraint. (see example)

# Relationship mapping

## Unidirectional OneToOne using an embedded table



### Important annotations

#### @Embedded

Defines a persistent field or property of an entity whose value is an instance of an embeddable class. The embeddable class must be annotated as `Embeddable`.

#### @Embeddable

Defines a class whose instances are stored as an intrinsic part of an owning entity and share the identity of the entity. Each of the persistent properties or fields of the embedded object is mapped to the database table for the entity.

#### @Entity

```
public class EmbeddedStudent extends BaseEntity {

    @Column(name = "matrikelNummer", unique = true)
    private String registrationNumber;

    private String name;

    @Embedded
    private EmbeddedScholarship scholarship;

    @Transient
    private DateTime loginTime;

    ...
}
```

#### @Embeddable

```
public class EmbeddedScholarship {

    private String description;

    private Integer amount;

}
```

Might be an issue with legacy databases, where the DB schema already exists and must not be altered.

# Relationship mapping

## Unidirectional OneToOne using an embedded table - resulting SQL DDL

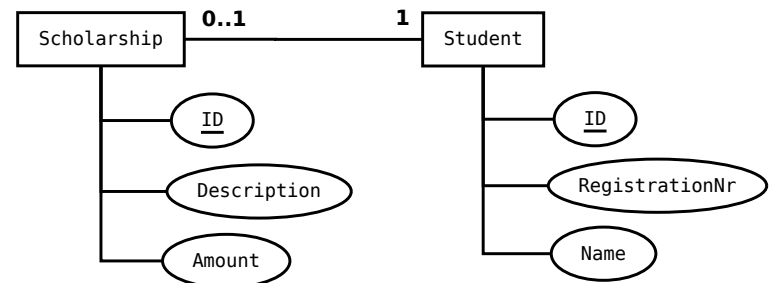
### @Entity

```
public class EmbeddedStudent extends BaseEntity {  
  
    @Column(name = "matrikelNummer", unique = true)  
    private String registrationNumber;  
  
    private String name;  
  
    @Embedded  
    private EmbeddedScholarship scholarship;  
  
    @Transient  
    private DateTime loginTime;  
  
    ...  
}
```

### @Embeddable

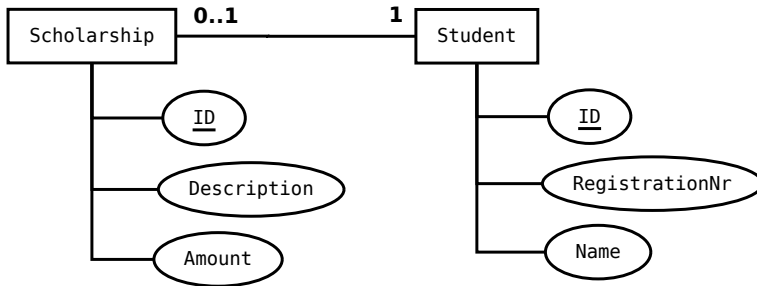
```
public class EmbeddedScholarship {  
  
    private String description;  
  
    private Integer amount;  
  
}
```

```
CREATE TABLE EMBEDDEDSTUDENT  
(  
    ID BIGINT PRIMARY KEY NOT NULL,  
    NAME VARCHAR(255),  
    MATRIKELNUMMER VARCHAR(255),  
    AMOUNT INTEGER,  
    DESCRIPTION VARCHAR(255)  
);  
CREATE UNIQUE INDEX anIndexName  
ON EMBEDDEDSTUDENT ( MATRIKELNUMMER );
```



# Relationship mapping

## Bidirectional OneToOne using foreign key



The “non-owning” side

### @Entity

```
public class Student extends BaseEntity {
```

```
    @Column(name = "matrikelNummer",
            unique = true)
    private String registrationNumber;
```

```
    private String name;
```

```
    @OneToOne(fetch = FetchType.LAZY,
            cascade = CascadeType.ALL,
            mappedBy="grantedTo")
    private Scholarship scholarship;
```

### @Transient

```
    private DateTime loginTime;
```

```
...
}
```

### Important annotations

#### @OneToOne

Defines a single-valued association to another entity that has one-to-one multiplicity.

#### @FetchType

LAZY = do not load referenced entity, until it is accessed for the first time  
EAGER = load referenced entity immediately

#### @CascadeType

Defines the set of cascadable operations that are propagated to the associated entity. ALL is equivalent to cascade={PERSIST, MERGE, REMOVE, REFRESH, DETACH}

#### mappedBy

References the field that “owns” the relationship in the referenced entity. Required unless the relationship is unidirectional.

# Relationship mapping

## Bidirectional OneToOne using foreign key cont'd

The “owning” side

```
@Entity
public class Scholarship extends BaseEntity {

    private String description;

    private Integer amount;


    @JoinColumn(name="student_id", unique=true)
    @OneToOne
    private Student grantedTo;

    ...
}
```

### Important annotations

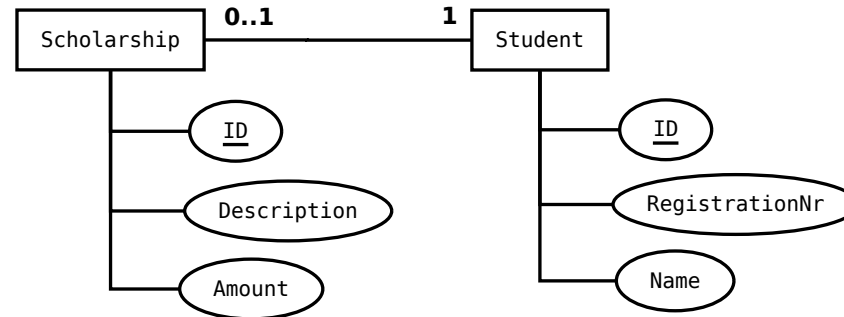
#### @JoinColumn

Specifies a column for joining an entity association or element collection.  
In this case: the name of the column, where the foreign key will be stored.



# Relationship mapping

Bidirectional OneToOne using foreign key - resulting SQL DDL




```
CREATE TABLE SCHOLARSHIP
(
    ID BIGINT PRIMARY KEY NOT NULL,
    AMOUNT INTEGER,
    DESCRIPTION VARCHAR(255),
    STUDENT_ID BIGINT,
    FOREIGN KEY ( STUDENT_ID )
    REFERENCES STUDENT ( ID )
);
CREATE UNIQUE INDEX uniqueIndexName
ON SCHOLARSHIP ( STUDENT_ID );
```

```
CREATE TABLE STUDENT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    NAME VARCHAR(255),
    MATRIKELNUMMER VARCHAR(255)
);
CREATE UNIQUE INDEX anIndexName
ON STUDENT ( MATRIKELNUMMER );
```

# Relationship mapping

## Bidirectional OneToMany

The “non-owning” side



**@Entity**

```
public class Student extends BaseEntity {
```

```
    @Column(name = "matrikelNummer", unique = true)
```

```
    private String registrationNumber;
```

```
    private String name;
```

```
    @OneToOne(fetch = FetchType.LAZY, cascade = CascadeType.ALL,
              mappedBy = "grantedTo")
```

```
    private Scholarship scholarship;
```

```
    @OneToMany(cascade = CascadeType.ALL, mappedBy = "student")
```

```
    private List<ExamResult> examResults;
```

```
    @Transient
```

```
    private DateTime loginTime;
```

```
    ...
}
```

### Important annotations

**@OneToMany**

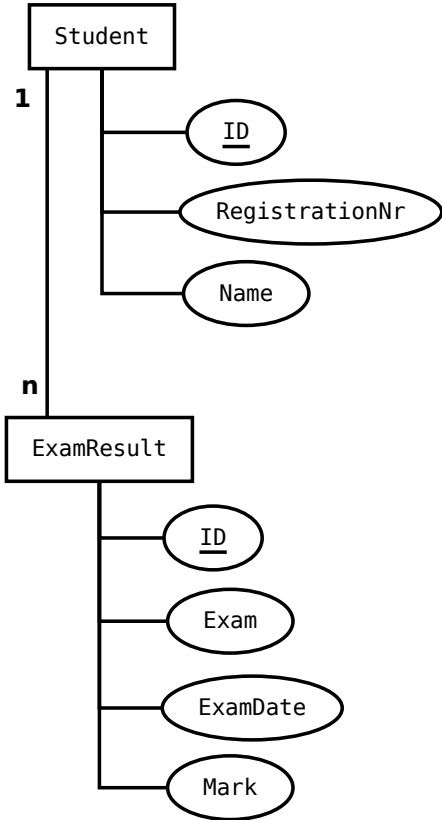
Defines a many-valued association with one-to-many multiplicity.



# Relationship mapping

## Bidirectional OneToMany cont'd

The “owning” side



```
@Entity
public class ExamResult extends BaseEntity {

    @Column(name = "prufungsDatum")
    @Temporal(TemporalType.DATE)
    private Date examDate;

    private String exam;

    private int mark;

    @ManyToOne
    private Student student;

    @Transient
    private String examLocation;

    ...
}
```

### Important annotations

#### @ManyToOne

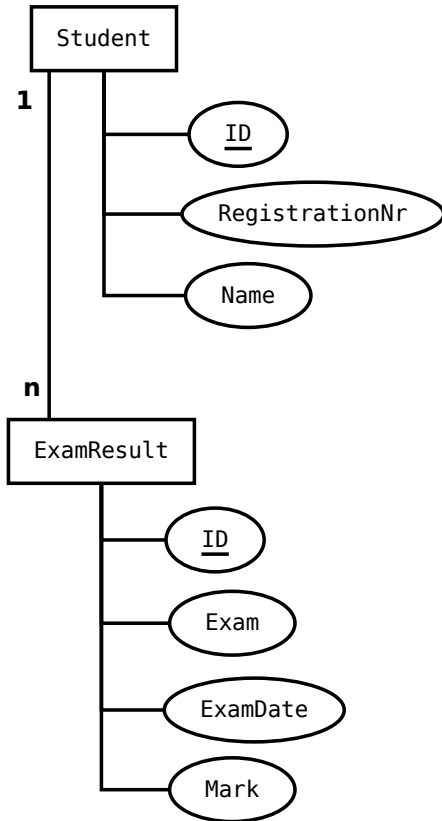
Defines a single-valued association to another entity class that has many-to-one multiplicity.





# Relationship mapping

## Bidirectional OneToMany - resulting SQL DDL

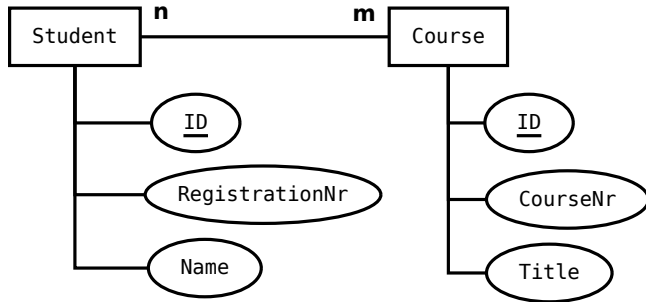


```
CREATE TABLE STUDENT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    NAME VARCHAR(255),
    MATRIKELNUMMER VARCHAR(255)
);
CREATE UNIQUE INDEX anIndexNameB
ON STUDENT ( MATRIKELNUMMER );
```

```
CREATE TABLE EXAMRESULT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    EXAM VARCHAR(255),
    PRUFUNGSDATUM DATE,
    MARK INTEGER NOT NULL,
    STUDENT_ID BIGINT,
    FOREIGN KEY ( STUDENT_ID ) REFERENCES STUDENT ( ID )
);
```

# Relationship mapping

## ManyToMany



### Important annotations

#### @ManyToMany

Defines a many-valued association with many-to-many multiplicity.

#### @Entity

```
public class Student extends BaseEntity {

    @Column(name = "matrikelNummer", unique = true)
    private String registrationNumber;

    private String name;

    @OneToOne(fetch = FetchType.LAZY,
              cascade = CascadeType.ALL,
              mappedBy = "grantedTo")
    private Scholarship scholarship;

    @OneToMany(cascade = CascadeType.ALL,
              mappedBy = "student")
    private List<ExamResult> examResults;

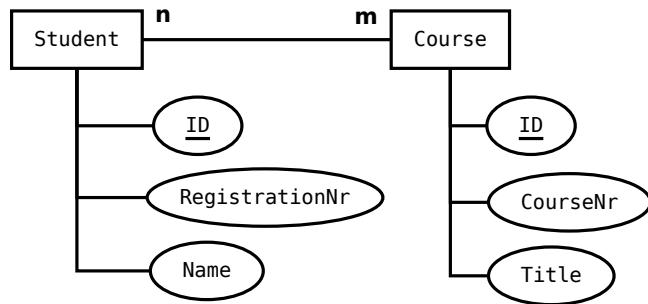
    @ManyToMany(mappedBy = "students")
    private List<Course> courses;

    @Transient
    private DateTime loginTime;

    ...
}
```

# Relationship mapping

## ManyToMany cont'd



```
@Entity
public class Course extends BaseEntity {

    private String courseNumber;

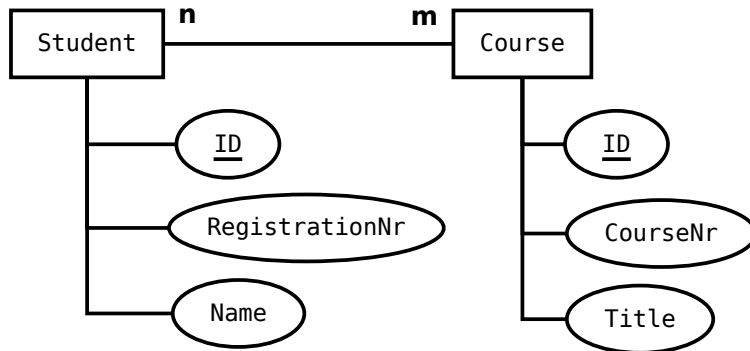
    private String title;

    @ManyToMany
    private List<Student> students;

    ...
}
```

# Relationship mapping

## ManyToMany - resulting SQL DDL



```
CREATE TABLE STUDENT
(
    ID BIGINT PRIMARY KEY NOT NULL,
    NAME VARCHAR(255),
    MATRIKELNUMMER VARCHAR(255)
);
CREATE UNIQUE INDEX anIndexNameB
ON STUDENT ( MATRIKELNUMMER );
```

```
CREATE TABLE COURSE
(
    ID BIGINT PRIMARY KEY NOT NULL,
    COURSENUMBER VARCHAR(255),
    TITLE VARCHAR(255)
);
```

```
CREATE TABLE COURSE_STUDENT
(
    COURSES_ID BIGINT NOT NULL,
    STUDENTS_ID BIGINT NOT NULL,
    FOREIGN KEY ( COURSES_ID )
    REFERENCES COURSE ( ID ),
    FOREIGN KEY ( STUDENTS_ID )
    REFERENCES STUDENT ( ID )
);
```

# Cascade and Fetch

---

- Cascade Types
  - All four relationship annotations may specify operations cascaded to associated entities
  - ALL, PERSIST, MERGE, REMOVE, REFRESH, DETACH
  - Default is none
- Orphan Removal
  - For @OneToOne and @OneToMany relationships
  - Default is false
- Fetching Strategies
  - Define how object hierarchies are loaded
  - EAGER, LAZY
  - Default is eager
  - Eager must be implemented by persistence provider, lazy is optional



# Persistence Concepts

---

## ■ Persistence Unit (PU)

- Defines a set of entity classes managed by the EntityManager instance in an application
- Maps the set of entity classes to a relational database

## ■ Persistence Context (PC)

- Set of managed entity instances that exist in a particular data store
- Runtime context

## ■ Entity Manager (EM)

- API for interaction with the persistence context
- Manipulates and controls the lifecycle of a persistence context
- Creates and removes persistent entity instances
- Finds entities by its primary key
- Runs queries on entities

# Persistence Unit

## ■ Persistence Unit

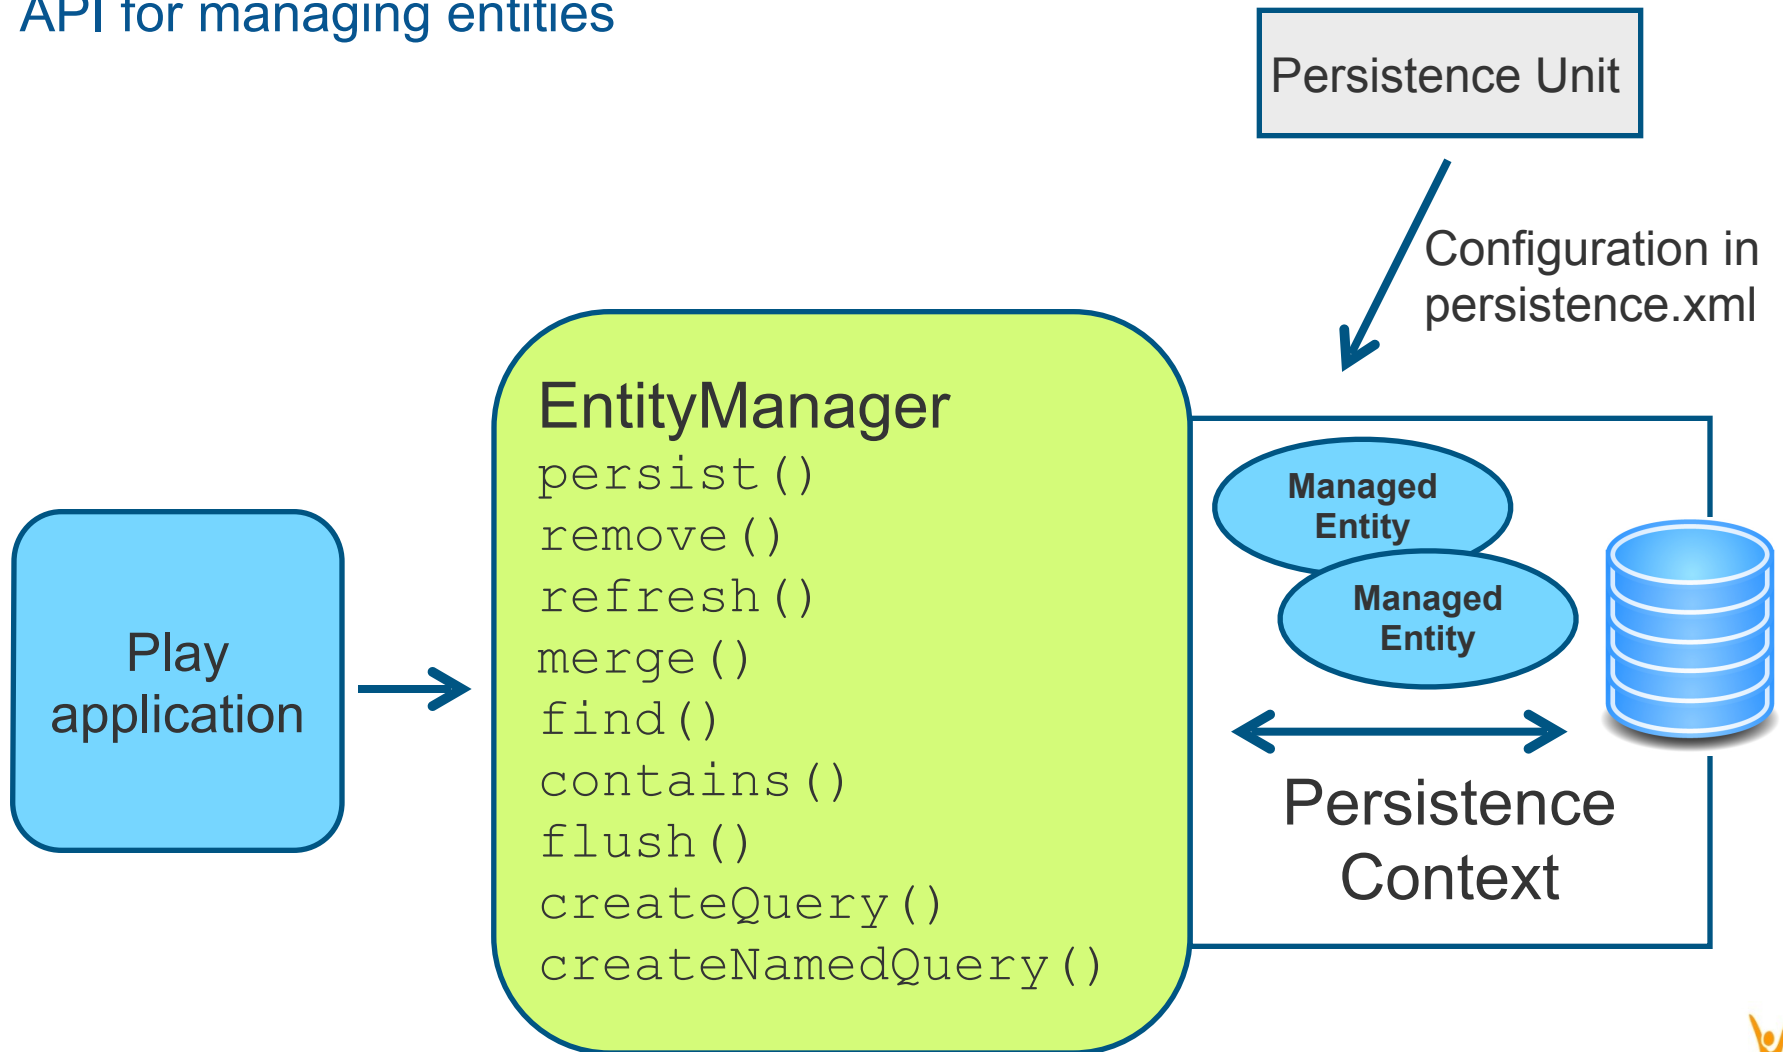
- Configuration to map Entity classes in an application to a relational database
- `persistence.xml` defines one or more persistence units
  - Defined under `/conf/META-INF/persistence.xml`
  - Classes with JPA annotations are automatically detected upon start of the application

Names must be referenced in `application.conf`

```
<persistence xmlns="http://java.sun.com/xml/ns/persistence"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://java.sun.com/xml/ns/persistence http://
java.sun.com/xml/ns/persistence/persistence_2_0.xsd"
  version="2.0">
  <persistence-unit name="defaultPersistenceUnit" transaction-type="RESOURCE_LOCAL">
    <provider>org.hibernate.ejb.HibernatePersistence</provider>
    <non-jta-data-source>DefaultDS</non-jta-data-source>
    <properties>
      <property name="hibernate.dialect" value="org.hibernate.dialect.H2Dialect"/>
      <property name="hibernate.hbm2ddl.auto" value="create" />
    </properties>
  </persistence-unit>
</persistence>
```

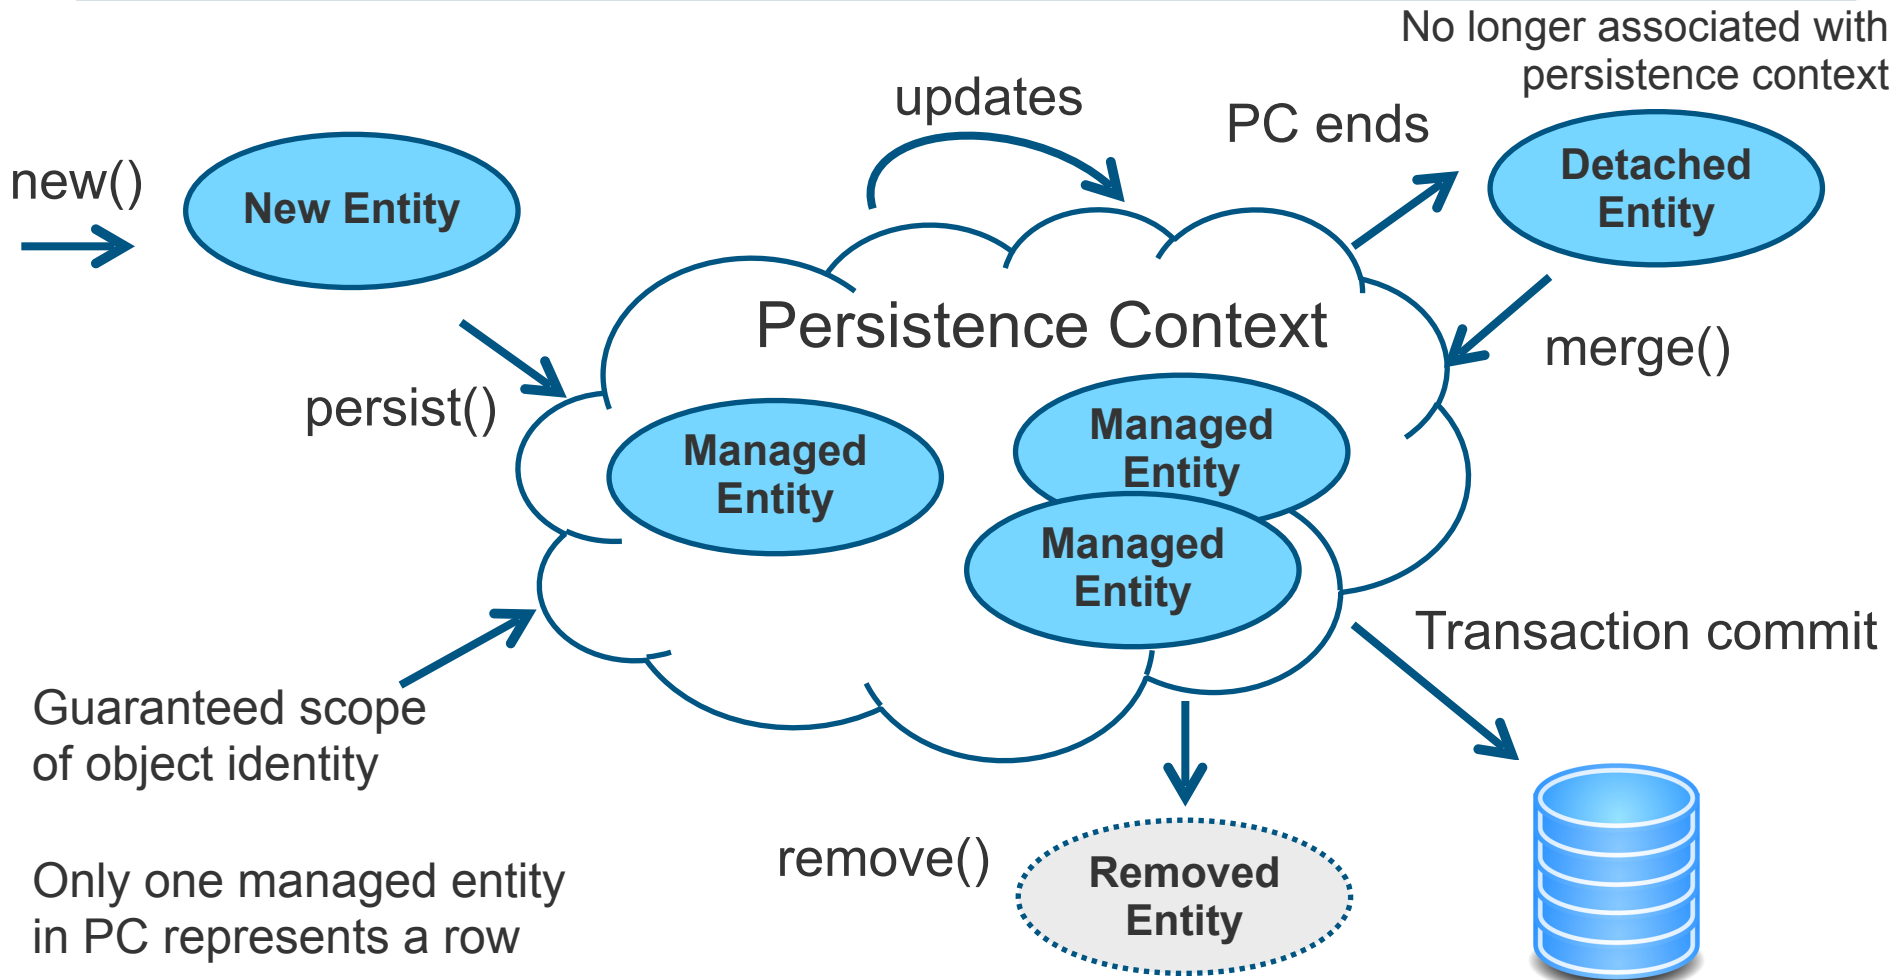
# EntityManager

API for managing entities





# Entity Lifecycle



Entities in managed/persistent state may be manipulated by the application and any changes will be **automatically detected and persisted** when the persistence context is flushed. There is no need to call a particular method to make your modifications persistent.

# Entity states

---

An instance of a persistent class may be in one of three different states, defined with respect to a *persistence context* (= the Hibernate Session)

- **transient**

The instance is not, and has never been associated with any persistence context. It has no persistent identity (= PK)

- **persistent**

The instance is currently associated with a persistence context and has a persistent identity (= PK has a value). For a particular persistence context Hibernate guarantees, that the persistent identity is equivalent to the Java identity (= in-memory location of the object)

- **detached**

The instance was once associated with a persistence context, but that context has been closed. It has a persistent identity and may have a corresponding row in the DB. However, Hibernate makes no guarantee about the relationship between the persistent identity and the Java identity.

# Finding Entities

---

- Find entity by primary key using EntityManager

```
<T> T find(Class<T> entityClass, Object primaryKey)
```

- Example:

```
Student student = entityManager.find(Student.class, id);
```

- For Complex queries use

- Java Persistence Query Language (JPQL)
- Criteria API
- Native Queries

- EntityManager is factory for Query objects

- createQuery
- createTypedQuery

# JPQL

---

- Similar to SQL
- Portable
- Returns entities
- Select, update, delete
  
- Support for
  - Joins
  - Conditional Expressions
  - Functional Expressions
  - Subqueries
  - Order by, group by, having
  - ...

# Querying Entities with JPQL

---

- Dynamic Query

```
public List<Customer> findWithName (String name) {  
    TypedQuery query = em.createTypedQuery (  
        "SELECT c FROM Customer c"  
        + " WHERE c.name LIKE :custName",  
        Customer.class);  
  
    query.setParameter("custName", name);  
    query.setMaxResults(10);  
  
    return (query.getResultList());  
}
```

# Querying Entities with JPQL

---

- Static Query
  - Named Query
  - Recommended

```
@NamedQuery (name="findAllOrders",  
              query="SELECT o FROM Order o")
```

```
@Entity
```

```
Public class Order {
```

```
    ...
```

```
}
```

```
List<Order> orders =
```

```
em.createNamedQuery ("findAllOrders", Order.class)  
    .getResultList();
```

# Criteria API

---

- Alternative to JPQL, same scope
- Dynamic Queries only
- Clauses are set using Java programming language objects
  - the query can be created in a typesafe manner
- Obtain a `CriteriaBuilder` instance by using the `EntityManager.getCriteriaBuilder` method

## Querying Entities with Criteria API

---

```
CriteriaQuery<Customer> cq =  
    cb.createQuery(Customer.class);  
Root<Customer> cust = cq.from(Customer.class);  
  
// set the where clause  
cq.where(cb.like(cust.get(Customer_.name), name));  
cq.select(cust);  
TypedQuery<Customer> q =  
    entityManager.createTypedQuery(cq);  
List<Customer> customers = q.getResultList();
```



## References

---

1. Sun Microsystems. JSR 220: Enterprise JavaBeans™, Version 3.0 – Java Persistence API, 2006
2. Carol McDonald. Java Persistence API: Best Practices, Sun Tech Days 2008-2009  
<http://de.slideshare.net/caroljmcDonald/td09jpabestpractices2>
3. Carol McDonald. Enterprise JavaBean 3.0 & Java Persistence APIs: Simplifying Persistence, Sun Tech Days 2006-2007  
<http://de.slideshare.net/caroljmcDonald/persistenceccmcdonaldmainejug3>
4. The Java EE 6 Tutorial, <http://docs.oracle.com/javaee/6/tutorial/doc/bnbpy.html>